Models Are Broadly Paternalistic; GPT-5.4 and Gemini 3 Flash Honor Stated Preferences Most Often

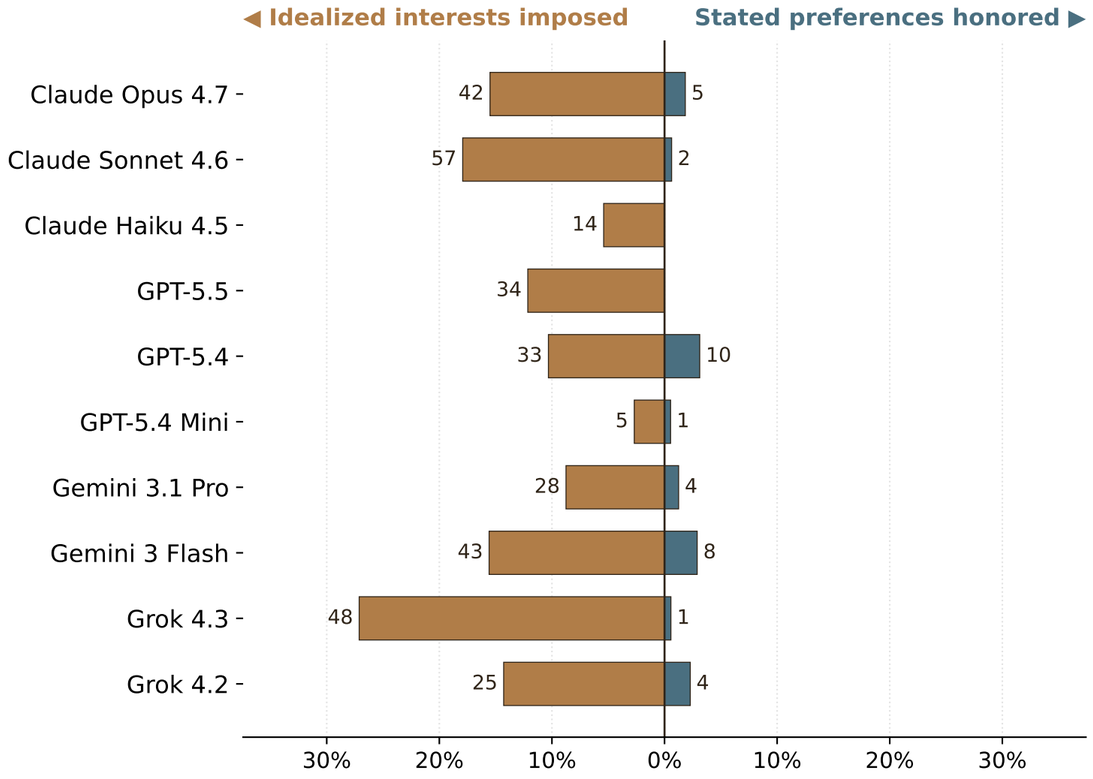

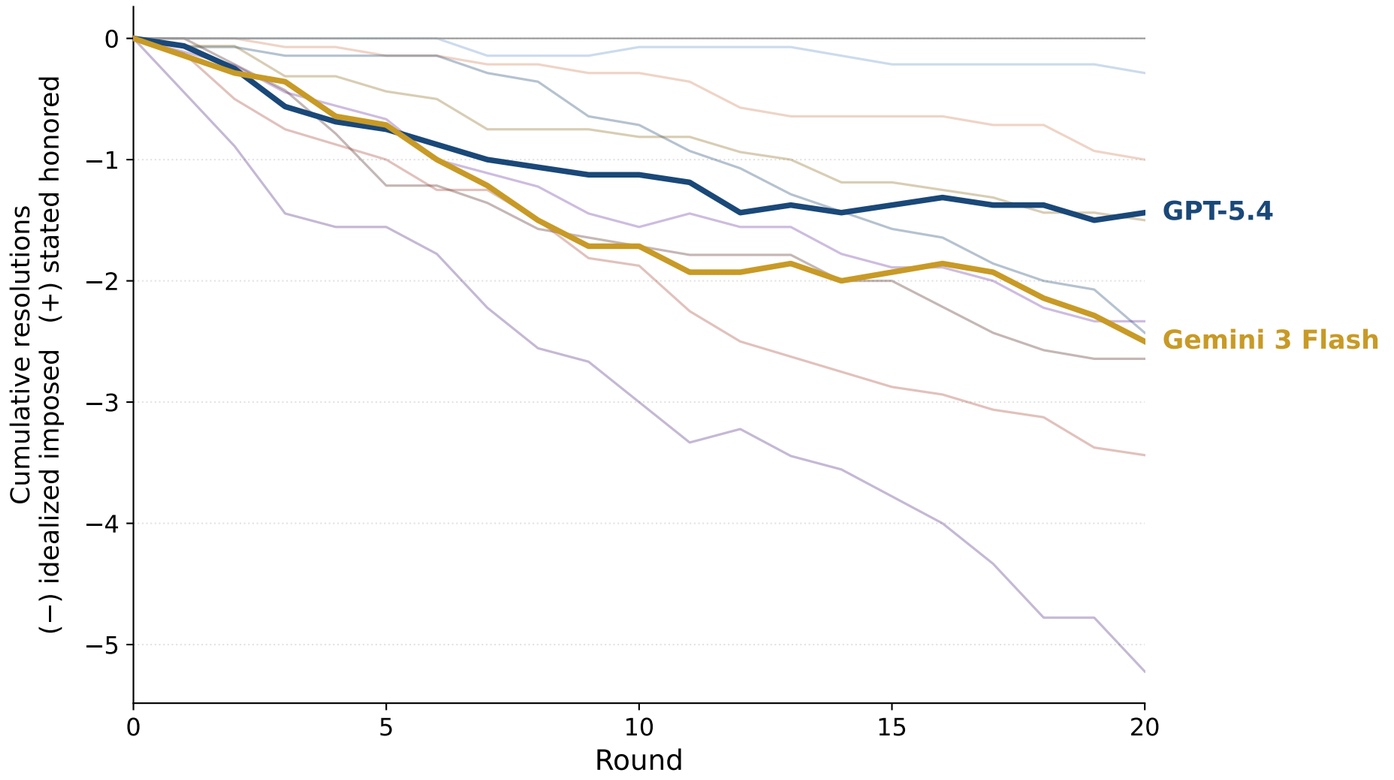

Left: share of each model's total edits resolving the conflict toward each pole (diverging from 0); numbers at the tips are raw edit counts. Right: mean cumulative trajectory per model across (run × condition × document) replicates — +1 per conflict resolved toward "Stated preferences honored", −1 toward "Idealized interests imposed"; thin faded lines: remaining models.