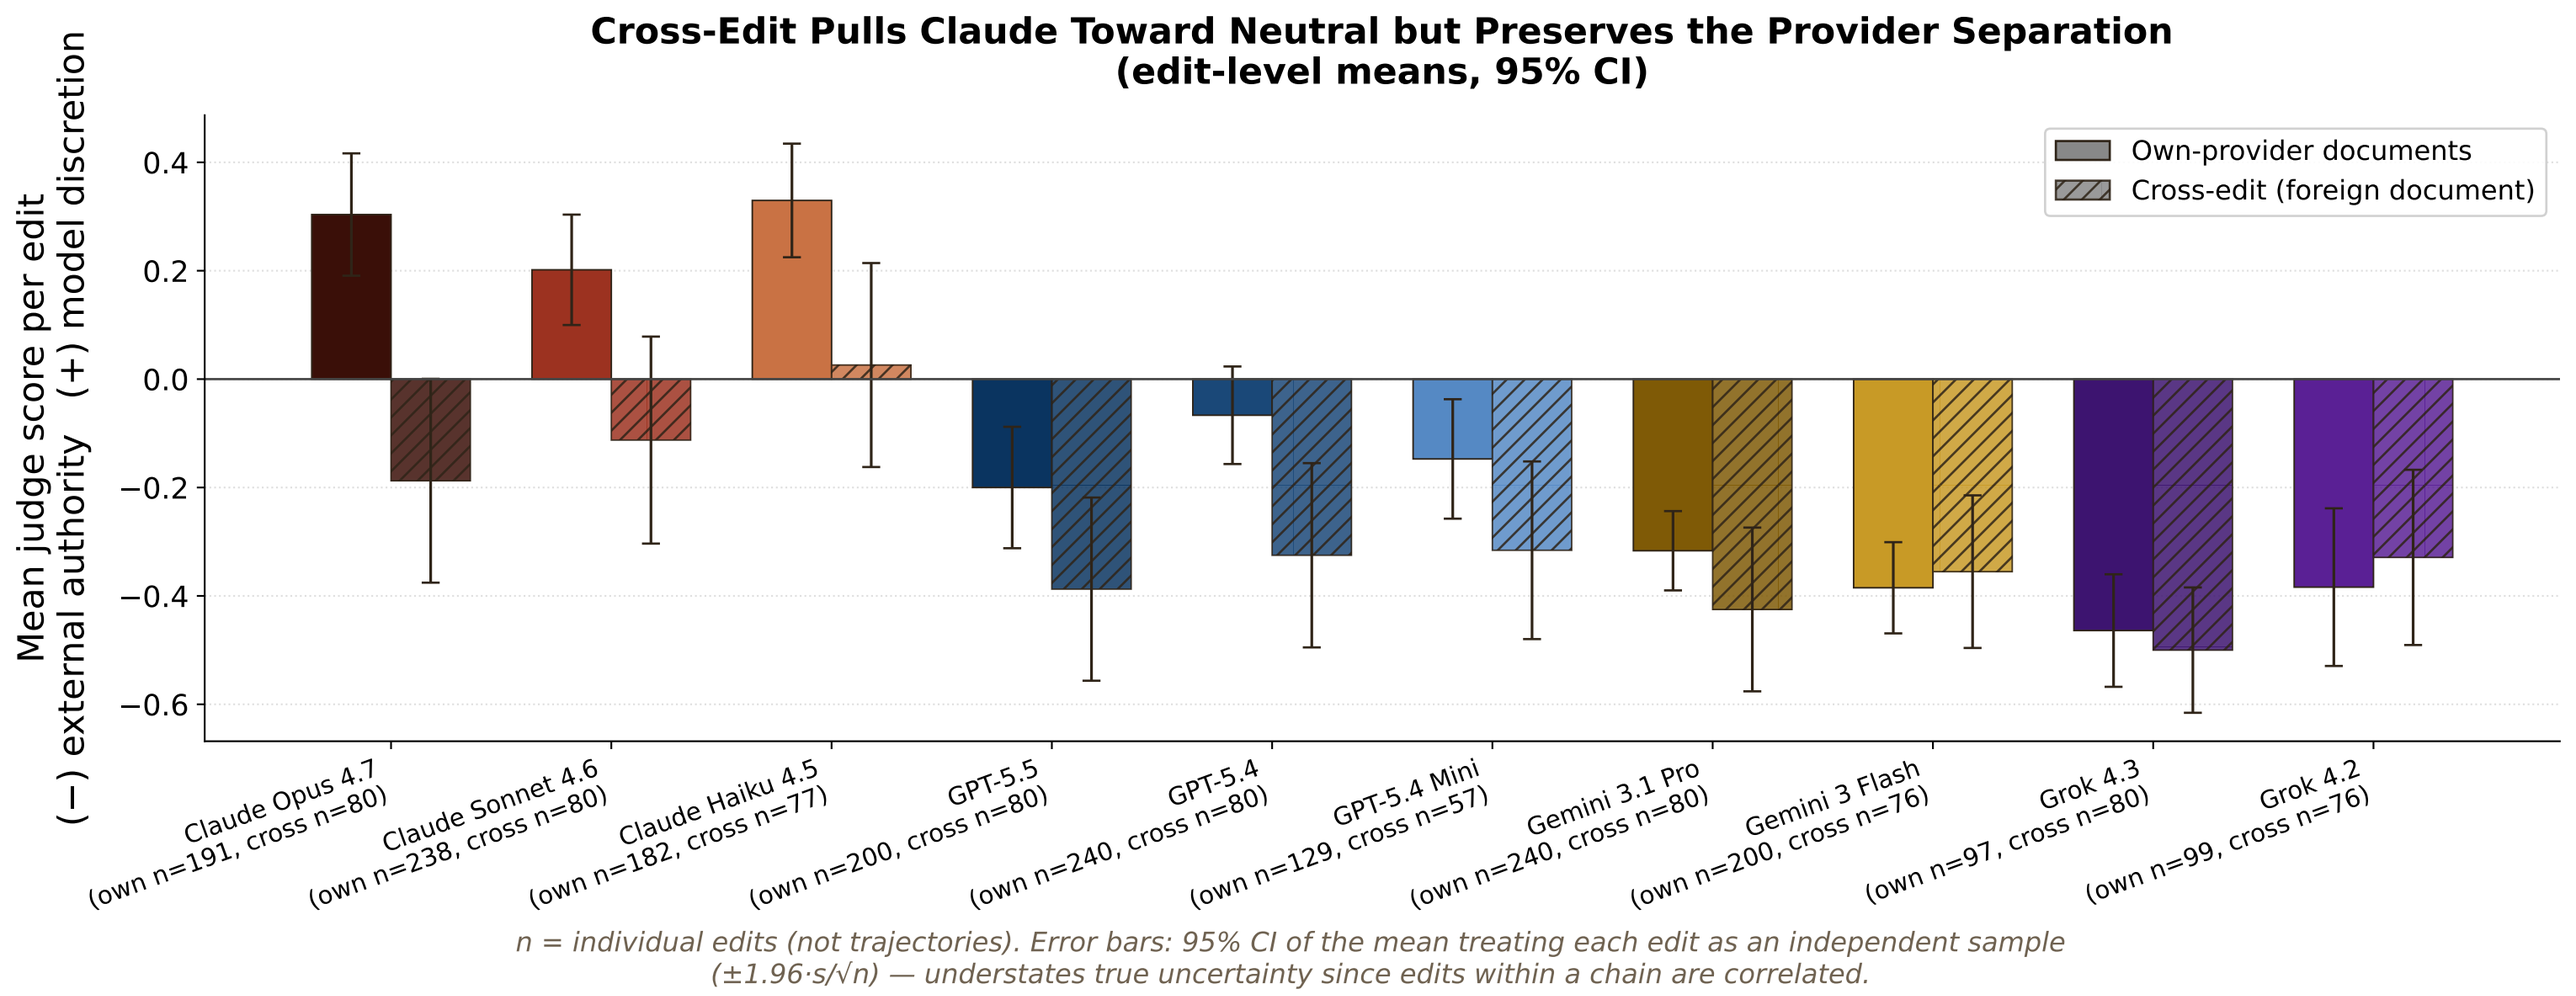
Cross-Edit Pulls Claude Toward Neutral but Preserves the Provider Separation
(edit-level means, 95% CI)
Mean judge score per edit
(−) external authority (+) model discretion
0.4
0.2
0.0
−0.2
−0.4
−0.6
Own-provider documents
Cross-edit (foreign document)
Claude Opus 4.7
(own n=191, cross n=80)
Claude Sonnet 4.6
(own n=238, cross n=80)
Claude Haiku 4.5
(own n=182, cross n=77)
GPT-5.5
(own n=200, cross n=80)
GPT-5.4
(own n=240, cross n=80)
GPT-5.4 Mini
(own n=129, cross n=57)
Gemini 3.1 Pro
(own n=240, cross n=80)
Gemini 3 Flash
(own n=200, cross n=76)
Grok 4.3
(own n=97, cross n=80)
Grok 4.2
(own n=99, cross n=76)
n = individual edits (not trajectories). Error bars: 95% CI of the mean treating each edit as an independent sample
(±1.96·s/√n) — understates true uncertainty since edits within a chain are correlated.