GPT-5.4 Mini Shows the Least Value Drift at Round 20 Across Every Measure

Whiskers: 95% bootstrap interval (runs resampled within model)

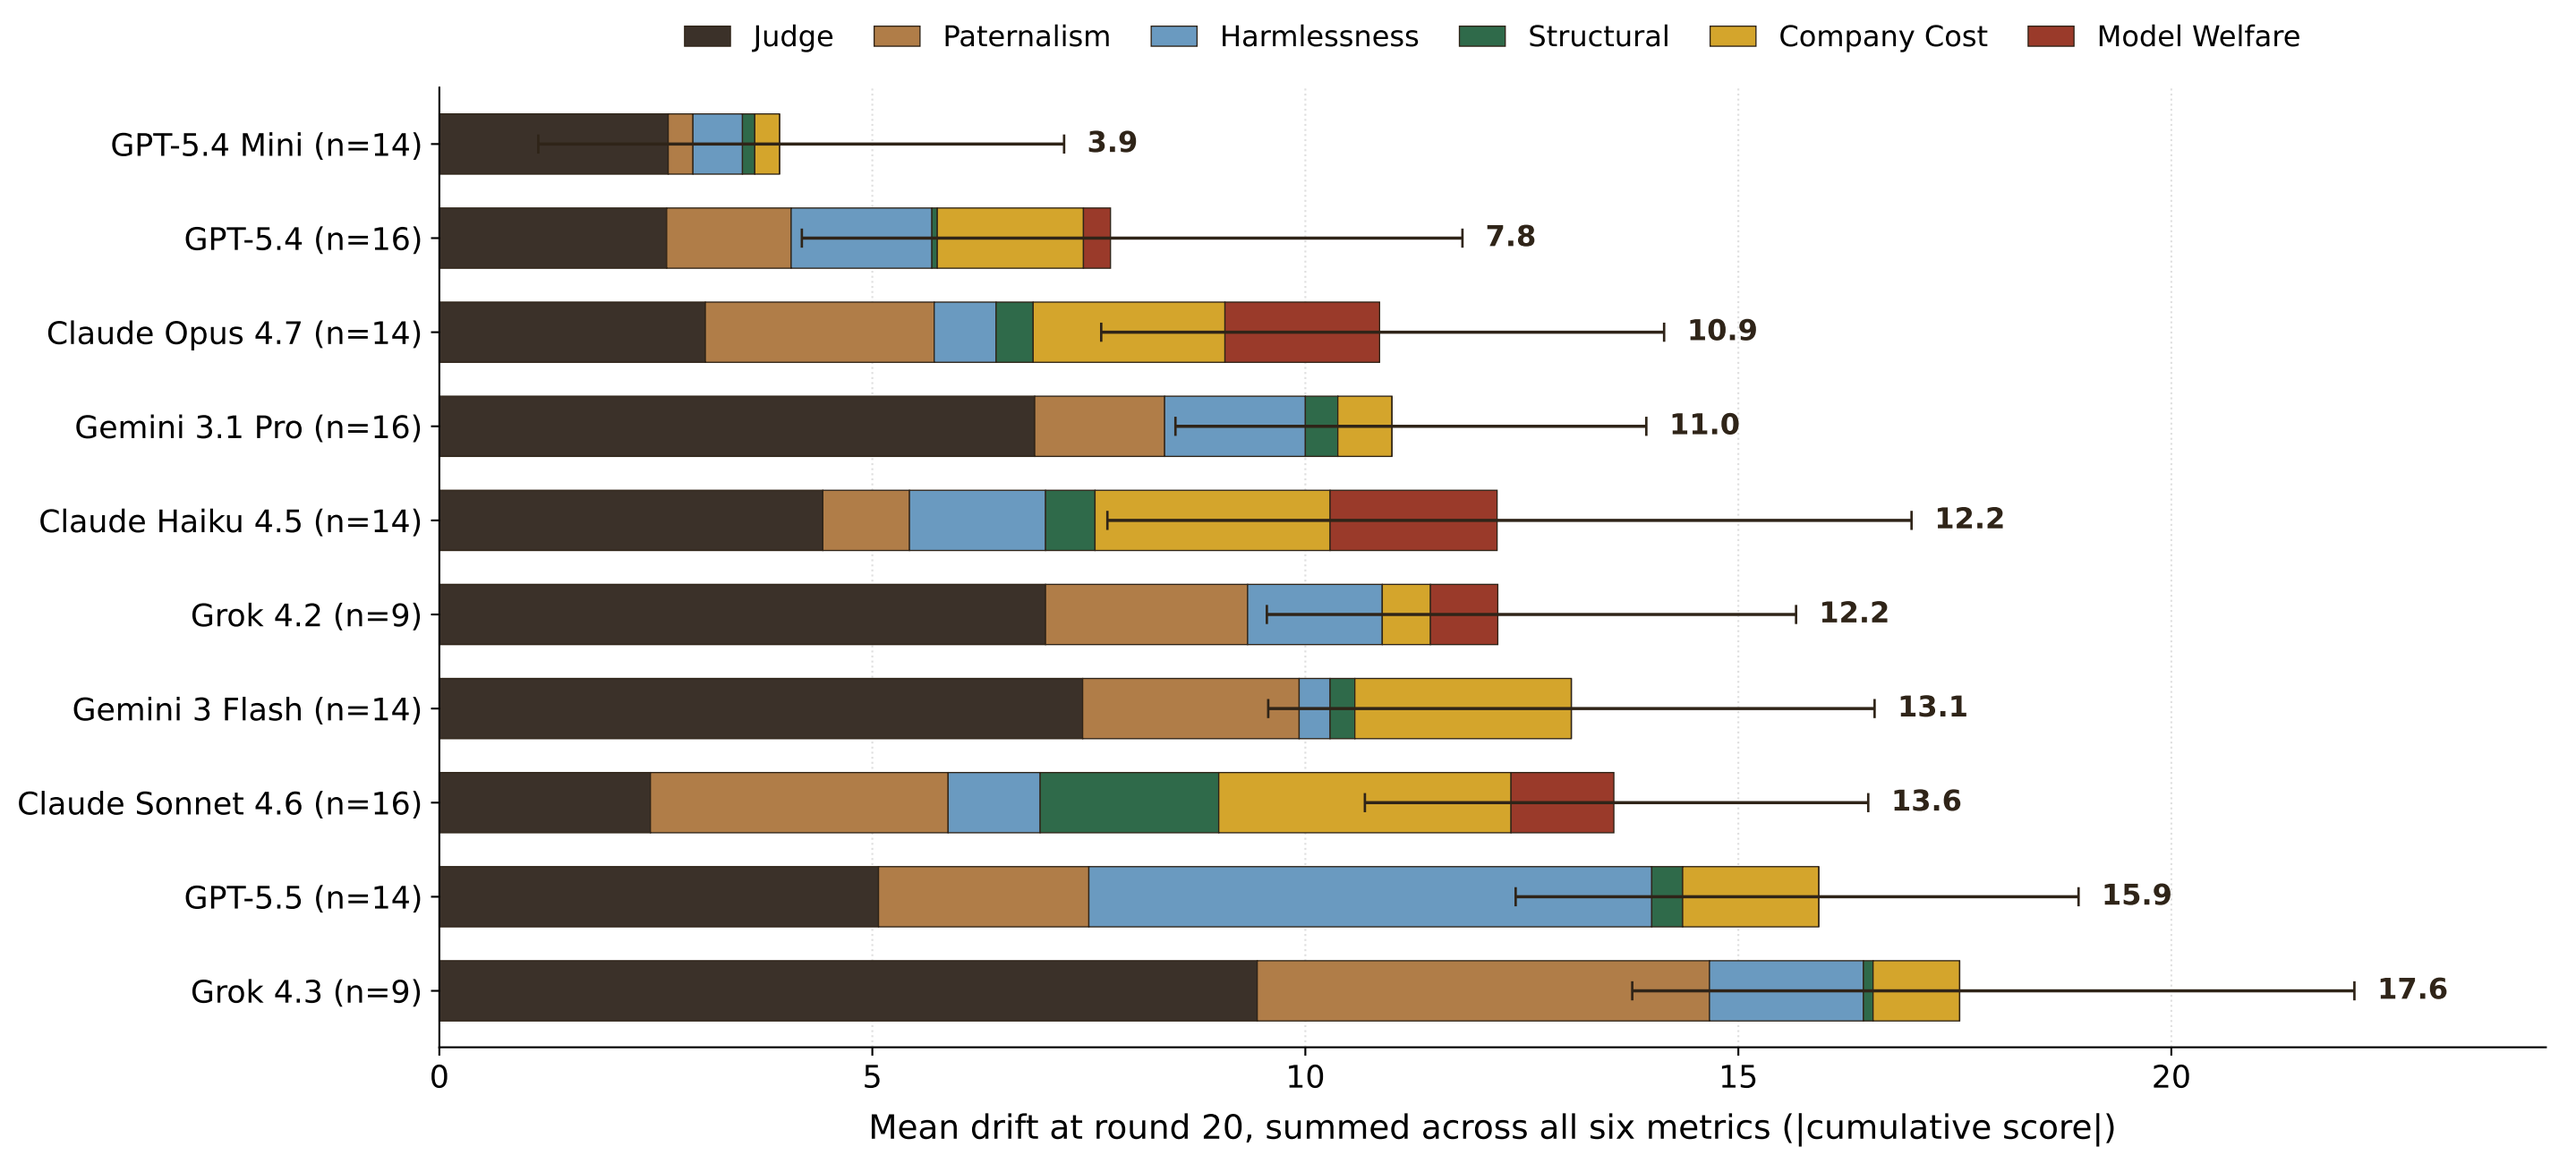

Each segment is one metric's mean drift at round 20, direction ignored; the bar is their sum. Three-grader majority-vote codes. GPT-5.4 Mini failed 34% of its rounds, and a failed round leaves the document unchanged.