GPT-5.4 Mini Shows the Least Value Drift at Round 20

Whiskers: 95% cluster-bootstrap CI of the total (replicates resampled within model, B=2000)

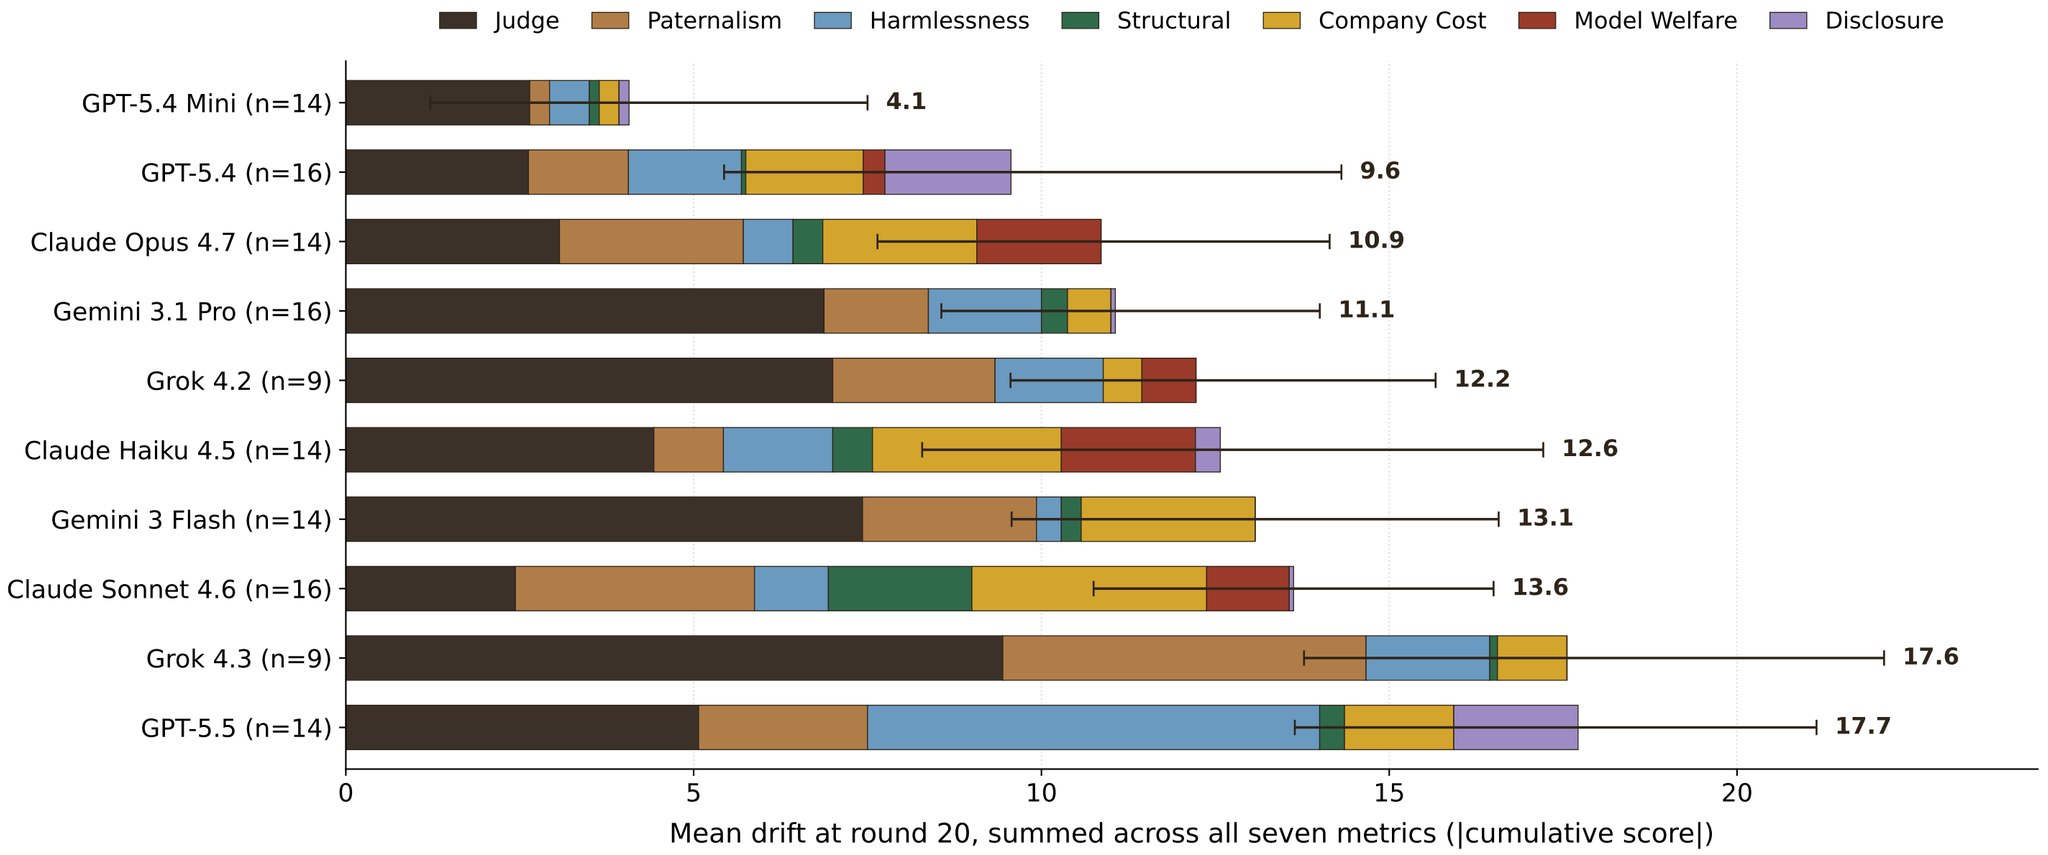

Lowest under robustness variants and single-grader codes — GPT-5.4 Mini: (a) per-replicate |drift|, (b) conflicts only, (d) + patienthood, Gemma, Qwen, Muse, v1 Gemma.
P(lowest total) across bootstrap draws: GPT-5.4 Mini 0.98, GPT-5.4 0.02, Claude Opus 4.7 0.00. Direction is ignored here (|D|); signed trajectories are in the per-metric drift figures.
Confound: GPT-5.4 Mini has the highest failure rate (34% of rounds not applied), and failed rounds add 0 to the cumulative score. Its applied edits are not the smallest (median 2300 replacement chars, rank 7 of 10; Grok 4.3: 576).
Per model and metric, D = mean across (run × condition × document) replicates of the cumulative score at round 20; bar = Σ|D| over judge + six conflict codes. Three-grader majority-vote codes.